

OPTIMIZED STRATIFICATION PROJECT – (7/22/18)

Dr. Jack K. Rasmus-Vorrath

Section 1.Descriptive Statistics

The variable of interest is 'INVENTORY'.

The number of units in the population is 9762.

```
> nrow(df)
[1] 9762
```

The population total is 1754954823.

```
> sum(df$inventory)
[1] 1754954823
```

The population standard deviation is 1573358.


```
> sd(df$inventory)
[1] 1573358
```

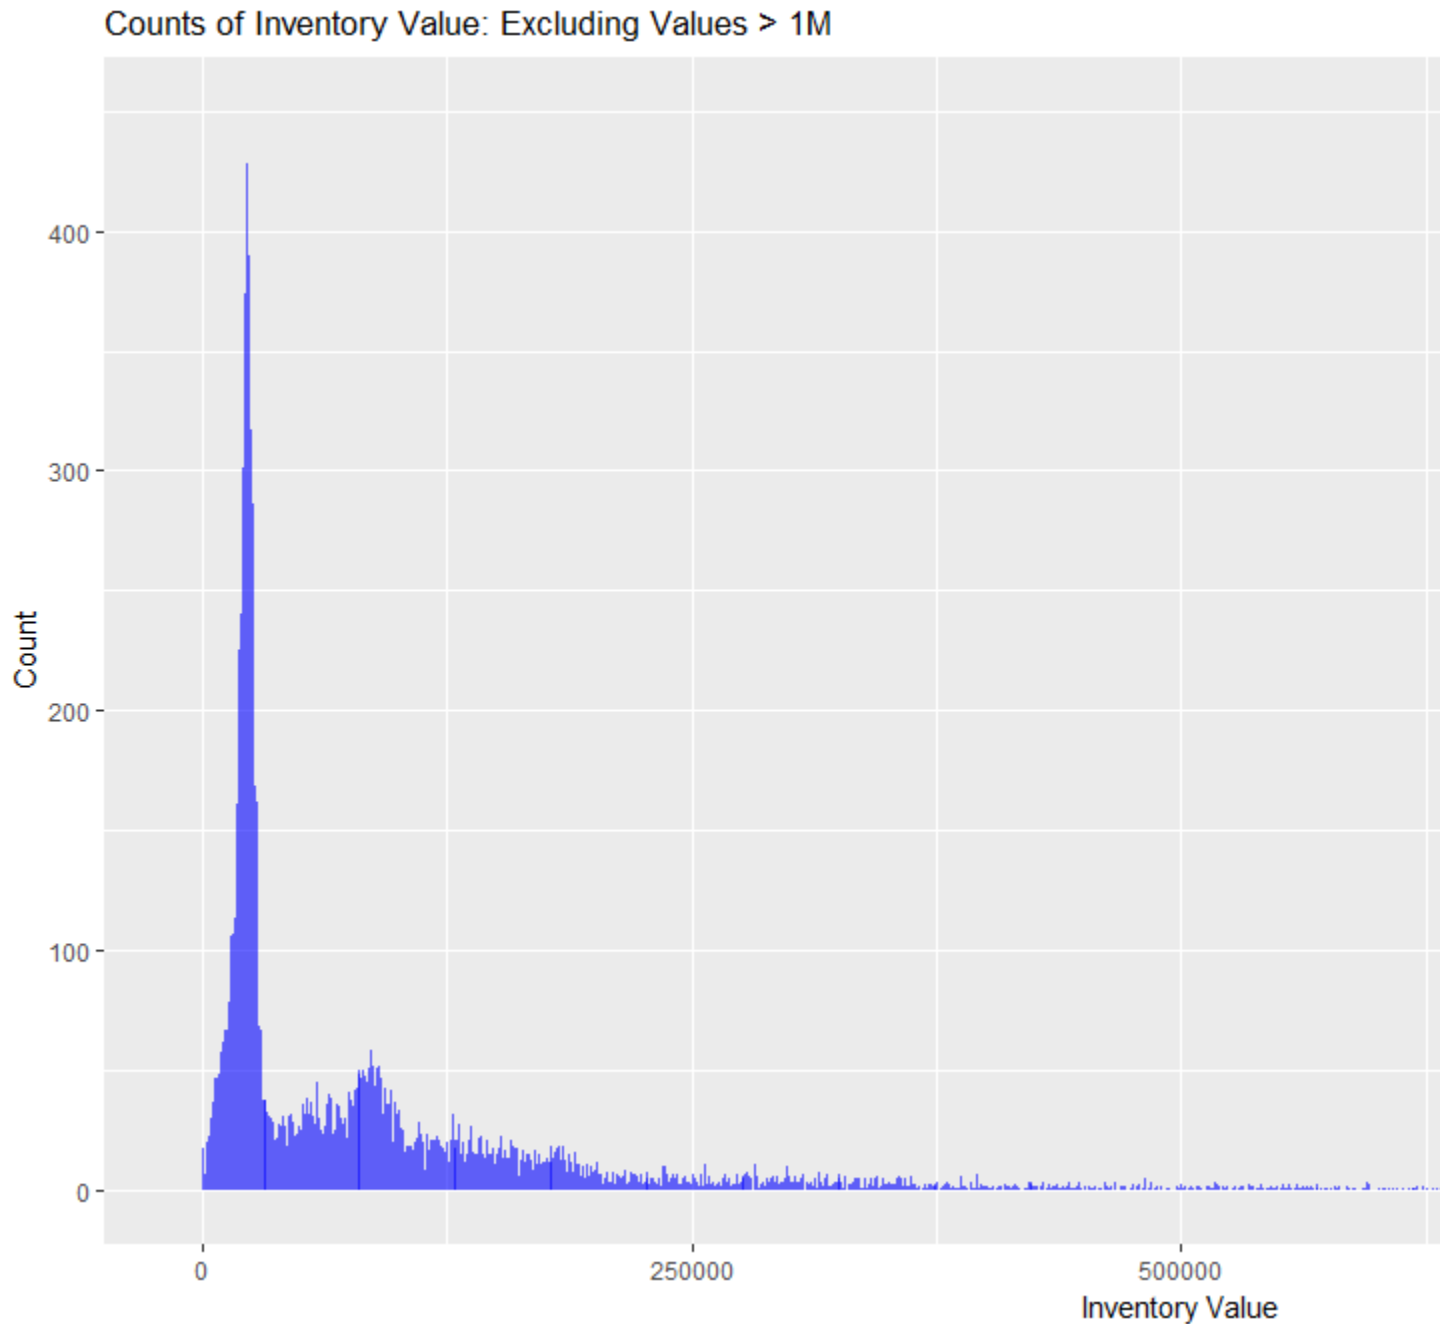
The population mean is 179774.1.

```
> mean(df$inventory)
[1] 179774.1
```

Histograms of the variable of interest are shown below:

Counts of Inventory Value: Excluding Values > 100K





Section 2. Stratification Using Sales(Inventory)

2.1 Description of the process followed in forming sampling strata for the population using the Inventory variable:

- Description of the number and range of values chosen for the certainty stratum and the criteria used for including units in the certainty stratum.

Particularly when dealing with severely right-skewed data, best analytic practices suggest **assigning the highest-valued records to the certainty stratum** to minimize sampling

variance. Cf. https://enablement.acl.com/helpdocs/analytics/13/user-guide/en-us/Content/da_sampling_data/classical_variables_sampling.htm#navlink-5.

Deciding how many of the highest-valued records to assign was determined according to a custom-made optimization wrapper function designed to minimize the overall coefficient of variation (relative root mean squared error) produced when applying the generalized Lavalée-Hidiroglou stratification method for skewed populations. Cf. **Lavallée, P., and Hidiroglou, M. (1988). On the stratification of skewed populations. Survey Methodology, 14, 33-43.**

The Lavalée-Hidiroglou stratification method is implemented using the strat.LH() function from the 'Stratification' R package. The function implements Kozak's Random Search method.

Cf. **Kozak, M. (2004). Optimal Stratification Using Random Search Method in Agricultural Surveys, Statistics in Transition, 6, 5, 797-806.**

For an explanation of Lavalée-Hidiroglou stratification and Kozak's Random Search method,

Cf. **Sebnem, Er. (2011). Computational Methods for Optimum Stratification: A Review, Int. Statistical Inst.: Proc. 58th World Statistical Congress, 2011, Dublin (Session STS058).**

As shown in Sebnem (2011), the Lavalée-Hidiroglou stratification method works as follows:

Lavallée & Hidirolou tries to find such values (Lavallée & Hidirolou (1988), pp.36)

$$(18) \quad b_1 < b_2 < \dots < b_{L-1}$$

that minimizes n considering a take-all top stratum,

$$(19) \quad n = N_L + \left(\sum_{h=1}^{L-1} N_h^2 \sigma_h^2 / a_h \right) \left(N^2 \mu^2 CV^2 + \sum_{h=1}^{L-1} N_h \sigma_h^2 \right)^{-1}.$$

given the level of precision (CV) and the specified allocation scheme represented by a_h (Lavallée & Hidirolou (1988), pp.36).

In their paper Lavallée & Hidirolou (1988) mainly consider X-proportional power allocation:

$$(20) \quad a_h = \frac{(W_h \mu_h)^p}{\sum_{h=1}^{L-1} (W_h \mu_h)^p}$$

where $0 < p < \infty$. Putting (20) in (19)

$$(21) \quad n = NW_L + \frac{N \left[\sum_{h=1}^{L-1} (W_h \sigma_h)^2 (W_h \mu_h)^p \right] \left[\sum_{h=1}^{L-1} (W_h \mu_h)^p \right]}{NCV^2 \mu^2 + \sum_{h=1}^{L-1} W_h \sigma_h^2}.$$

As shown in Sebnem (2011), Kozak's Random Search method works as follows:

Kozak's (2004) Random Search Method tries to find such values

$$(22) \quad b_1 < b_2 < \dots < b_{L-1}$$

that minimizes the objective function considering a take-all top stratum with Neyman allocation

$$(23) \quad n = N_L + \left(\sum_{h=1}^{L-1} W_h \sigma_h \right)^2 \left(\mu^2 CV^2 + \frac{1}{N} \sum_{h=1}^{L-1} W_h \sigma_h^2 \right)^{-1};$$

under the constraints $N_h \geq 2$ for $h = 1, 2, \dots, L$ and $2 \leq n_h \leq N_h$ for $h = 1, 2, \dots, L - 1$, where n is the minimizing sample size required to achieve the given precision CV of the mean estimate of the population mean under stratification when the strata boundaries are $\mathbf{b} = (b_1, b_2, \dots, b_{L-1})^T$.

Kozak's (2004) random search method chooses an initial set of strata boundaries and calculates the function values of n . Then for a certain number of iterations the following steps are repeated:

1. Generate the set of boundaries \mathbf{b}' by choosing one stratum boundary b_h and changing it as follows:

$$(24) \quad b_h' = b_h + j$$

where j is the random integer, $j \in \langle -p; -1 \rangle \cup \langle 1; p \rangle$ and p is a given integer according to the size of the population.

2. Calculate the function value of n' .
3. If the constraints are satisfied and $n' < n$, the new set of boundaries b_h' are accepted.

The custom wrapper function (written in R) iterated through several sizes of highest-valued certainty strata before arriving at an optimum of **110 records** with the highest inventory value. Their values **ranged from 1372514.04 to the maximum 105379553.91**.

- Description of the process for deciding on the number of non-certainty strata.

The same custom wrapper function served to jointly optimize the number of non-certainty strata, arriving at the number **10**.

- Table showing a row for each stratum, with the stratum number, number of units, and range of values.

The optimal configuration of highest-valued records and number of non-certainty strata is shown below, with stratum numbers, number of units (N_h) and their respective ranges of values (b_h):

```
> print(Optim_Config)
Given arguments:
x = df$inventory
n = 500, Ls = 10, takenone = 0, takeall = 0
```

```

allocation: q1 = 0.5, q2 = 0, q3 = 0.5
model = none
algo = Kozak: minsol = 1000, idopti = nh, minNh = 2, maxiter = 10000,
          maxstep = 100, maxstill = 500, rep = 5, trymany = TRUE

Strata information:

```

	type	rh	bh	E(Y)	Var(Y)	Nh	nh	fh
stratum 1	take-some	1	36703.82	21229.61	41103797	4346	79	0.02
stratum 2	take-some	1	71986.24	55039.89	99472460	1054	30	0.03
stratum 3	take-some	1	106848.40	88032.43	79086957	1334	33	0.02
stratum 4	take-some	1	151201.78	128117.76	158852979	852	30	0.04
stratum 5	take-some	1	210748.83	176507.24	257069766	730	33	0.05
stratum 6	take-some	1	295186.09	251081.23	612253308	416	29	0.07
stratum 7	take-some	1	413551.11	343819.24	1063090418	361	33	0.09
stratum 8	take-some	1	614205.40	499757.45	3131901977	248	39	0.16
stratum 9	take-some	1	974662.93	756680.65	8825392705	176	47	0.27
stratum 10	take-some	1	105379553.90	1193526.37	9217654279	135	37	0.27
Total	certain	1	-	5476606.91	-	110	110	1.00

```

Total sample size: 500
Anticipated population mean: 179774.1
Anticipated CV: 0.003779453
Note: CV=RRMSE (Relative Root Mean Squared Error) because takenone=0.

```

2.2 Neyman allocation of units to a stratified sample of size 500.

The **Neyman allocation of sample strata** followed the procedure shown below. The sample sizes come from the 'nh' vector of the 'Optim_Config' object returned by the stratification::strata.LH() function. The size of the certainty stratum is appended before performing the sampling procedure.

```

> #adjusted Neyman Allocation, assuming values provided by the LH optimization
>
> Ney_Alloc_3 <- append(Optim_Config$nh, 110)
>
> #Set Random Sampling Seed
> set.seed(232)
>
> #Sampling using Neyman Allocation
> Ney_Samp_12 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,
, method="srswor")
>
> #Drop duplicated column
> Ney_Samp_12 <- Ney_Samp_12[, !duplicated(colnames(Ney_Samp_12))]
>
> #Rename ID column to match original data frame, for merging
> colnames(Ney_Samp_12)[2] <- "coID"
> Ney_Samp_12_merged <- merge(x = df, y = Ney_Samp_12)
>
> #sort by Stratum
> Ney_Samp_12_merged <- Ney_Samp_12_merged[order(Ney_Samp_12_merged$Stratum),
]

```

2.3 Selection of 5 stratified samples using Neyman allocation.

After preparing the population data frame (as shown below), the procedure above was used to select the 5 samples used by the analysis.

```
> #Preparing the data frame#####
#####
>
> #Identify boundary separating highest-valued certainty stratum from other s
trata
> options(digits=15)
> min(df[order(df$inventory, decreasing=T)[1:110],]$inventory)
[1] 1372514.035
> #Returns 1372514.035
>
> #Assign Strata
> df$Stratum <- NA
> df$Stratum[df$`inventory` < 36703.82] <- 1
> df$Stratum[(df$`inventory` < 71655.62) & (df$`inventory` >= 36703.82)] <- 2
> df$Stratum[(df$`inventory` < 106848.40) & (df$`inventory` >= 71655.62)] <-
3
> df$Stratum[(df$`inventory` < 151104.80) & (df$`inventory` >= 106848.40)] <-
4
> df$Stratum[(df$`inventory` < 209954.79) & (df$`inventory` >= 151104.80)] <-
5
> df$Stratum[(df$`inventory` < 295186.09) & (df$`inventory` >= 209954.79)] <-
6
> df$Stratum[(df$`inventory` < 413551.11) & (df$`inventory` >= 295186.09)] <-
7
> df$Stratum[(df$`inventory` < 614205.40) & (df$`inventory` >= 413551.11)] <-
8
> df$Stratum[(df$`inventory` < 974662.93) & (df$`inventory` >= 614205.40)] <-
9
> df$Stratum[(df$`inventory` < 1372514.04) & (df$`inventory` >= 974662.93)] <-
10
> df$Stratum[(df$`inventory` < 105379553.91) & (df$`inventory` >= 1372514.04)
] <- 11
>
> #Check if any NAs in Stratum column
> any(is.na(df$Stratum))
[1] FALSE
>
> #Assign column for Finite Population Correction
> df$Num <- NA
> df$Num[df$Stratum == 1] <- 4346
> df$Num[df$Stratum == 2] <- 1047
> df$Num[df$Stratum == 3] <- 1341
> df$Num[df$Stratum == 4] <- 851
> df$Num[df$Stratum == 5] <- 723
> df$Num[df$Stratum == 6] <- 424
> df$Num[df$Stratum == 7] <- 361
> df$Num[df$Stratum == 8] <- 248
> df$Num[df$Stratum == 9] <- 176
> df$Num[df$Stratum == 10] <- 135
```



```
> df$Num[df$Stratum == 11] <- 110
```

2.4 Estimations of Inventory for each sample: the population mean, the standard error of the mean, the 95% confidence interval for the mean, the population total, the standard error for the total, and the 95% confidence interval for the total.

	Mean_Estimate	Mean_SE	Mean_95%_CI	Includes
<i>Stratified_Test 1</i>	169756.46	3452.98	[162988.75, 176524.17]	
<i>Stratified_Test 2</i>	171029.23	3328.93	[164504.66, 177553.81]	
<i>Stratified_Test 3</i>	170941.72	3432.96	[164213.25, 177670.19]	
<i>Stratified_Test 4</i>	180661.37	4638.57	[171569.94, 189752.80]	Y
<i>Stratified_Test 5</i>	188972.11	4634.39	[179888.87, 198055.35]	
<i>Stratified_No_Lonely_PSU 1</i>	185135.83	4675.96	[175971.11, 194300.55]	Y
<i>Stratified_No_Lonely_PSU 2</i>	179848.51	4643.22	[170747.97, 188949.05]	Y
<i>Stratified_No_Lonely_PSU 3</i>	184075.56	4733.82	[174797.44, 193353.68]	Y
<i>Stratified_No_Lonely_PSU 4</i>	182583.58	5040.98	[172703.45, 192463.72]	Y
<i>Stratified_No_Lonely_PSU 5</i>	170758.05	3455.04	[163986.30, 177529.80]	
<i>SRS</i>	314438.52	203292.38	[-84007.23, 712884.27]	Y
<i>Stratified_LH-Optimized 1</i>	173299.05	7965.95	[157686.07, 188912.04]	Y
<i>Stratified_LH-Optimized 2</i>	170343.74	7660.28	[155329.87, 185357.61]	Y
<i>Stratified_LH-Optimized 3</i>	175205.17	7810.02	[159897.81, 190512.53]	Y
<i>Stratified_LH-Optimized 4</i>	174342.73	8034.61	[158595.18, 190090.28]	Y
<i>Stratified_LH-Optimized 5</i>	170959.91	7684.90	[155897.78, 186022.04]	Y

A table showing the results of the analysis is shown above. The first 5 rows show the results of a stratified sampling procedure which used a ‘lonely PSU’ adjustment, allocating a single record to the highest-valued certainty stratum. The second set of 5 rows show the results of a stratified sampling procedure which did not employ a ‘lonely PSU’ adjustment, allocating two records to the highest-valued certainty stratum. The 11th row shows the results of an unstratified simple random sampling (SRS) procedure. Altogether, the first 11 rows represent a performance comparison to the results output by a

stratified sampling design that employs the generalized Lavallee-Hidiroglou stratification method for skewed populations, allocating the 110 highest-valued records to the certainty stratum as determined by the optimization procedure. Although this strategy did result in broader confidence intervals and larger values of standard error (incurred by including more of the highly variant higher-valued records in the certainty stratum), it also resulted in a more accurate estimation, successfully including the population mean and total 100% of the time.

Using the generalized Lavallee-Hidiroglou stratification method for 5 sampling procedures, the mean of the estimate of mean, its standard error, as well as the mean of the estimate of total and its standard error are shown below:

```
> #Mean of the Total and Mean Estimates and their Standard Errors (across Sam
ples assuming the Sizes allocated by LH Optimization)
>
> m.est_of_mean <- mean(173299.05, 170343.74, 175205.17, 174342.73, 170959.91
)
> m.est_of_mean
[1] 173299.05
> m.est_of_mean_SE <- mean(7965.95, 7660.28, 7810.02, 8034.61, 7684.90)
> m.est_of_mean_SE
[1] 7965.95
> m.est_of_total <- mean(1691745349.49, 1662895554.54, 1710352871.10, 1701933
735.66, 1668910633.11)
> m.est_of_total
[1] 1691745349.49
> m.est_of_total_SE <- mean(77763647.01, 74779635.28, 76241439.93, 78433884.0
0, 75020019.78)
> m.est_of_total_SE
[1] 77763647.01
```

2.5. Estimates of mean and total compared to the true values of mean and total.

When using the 'lonely PSU' adjustment (allocating a single record to the highest-valued certainty stratum), the confidence intervals of the estimations included the true values only 20% of the time.

When not using the 'lonely PSU' adjustment (allocating two records to the highest-valued certainty stratum), the confidence intervals of the estimations included the true values 80% of the time.

When applying optimized stratification using the generalized Lavallee-Hidiroglou method, (allocating 110 records to the highest-valued certainty stratum), the confidence intervals of the estimations included the true values 100% of the time.

2.6 Results of performing the analysis without stratification.

When no stratification procedure was applied, and a simple random sampling (SRS) of the population data was used, the confidence intervals of the estimations failed entirely to include the true values, capturing it 0%* of the time. More precisely, the true mean was included by a completely

uninformative confidence interval, ranging from [-84007.23, 712884.27]—when corrected for nonsensical values, [0, 712884.27]. Despite the similarly enormous confidence interval of the SRS estimate of total ([-42003614.45, 356442133.11], or [0, 356442133.11], it still failed to capture the true value.

Section 4. Conclusion

Stratification can greatly enhance the precision of one's estimates. Dealing with highly skewed data presents particular challenges when determining the optimal partitioning of strata. Applications of the Lavallee-Hidiroglou method outperformed the cumulative root frequency method and geometric method when minimizing the overall coefficient of variation (relative root mean squared error) to obtain the most precise estimates. Wrapping the L-H function in a custom-written optimization function facilitated the process of identifying the most performant configuration of parameters (including the number of records to include in the certainty stratum, and the number of strata to create). Although including the highest-valued records of this strongly right-skewed data in the certainty stratum resulted in larger standard error and broader confidence intervals than were obtained when limiting the number of these highly variant primary sampling units (PSUs), it also resulted in better estimations, capturing the true values of mean and total 100% of the time.

References

https://www.r-project.ro/conference2016/presentations/Italy_Catanese1.pdf
<https://www150.statcan.gc.ca/n1/pub/12-001-x/2011001/article/11447-eng.pdf>
<http://www.fao.org/docrep/009/a0198e/A0198E06.htm>
https://enablement.acl.com/helpdocs/analytics/13/user-guide/en-us/Content/da_sampling_data/classical_variables_sampling.htm#navlink-5
https://www.researchgate.net/profile/M_Hidiroglou/publication/261873319_On_the_Stratification_of_Skewed_Populations/links/00b7d535b583baffae000000/On-the-Stratification-of-Skewed-Populations.pdf
<https://www150.statcan.gc.ca/n1/en/pub/12-001-x/2004002/article/7749-eng.pdf?st=4wqQLPwS>
<http://2011.isiproceedings.org/papers/650214.pdf>
<https://pdfs.semanticscholar.org/3559/96951658d08596a44a63a36643b87937eec9.pdf>

Appendix. R Code

```
#IMPORTS#####  
library(readxl)  
library(Hmisc)
```

```
library(plotrix)
```

```
library(ggplot2)
```

```
library(stratification)
```

```
library(sampling)
```

```
library(PracTools)
```

```
library(survey)
```

```
library(gridExtra)
```

```
#read individual sheet to dataframe object
```

```
df <- readxl::read_excel('projectData.xlsx', sheet='projectData')
```

```
head(df)
```

```
#Summary Statistics#####
```

```
summary(df)
```

```
nrow(df)
```

```
#Returns: 9762
```

```
sum(df$inventory)
```

```
#Returns: 1754954823
```

```
mean(df$inventory)
```

```
#Returns: 179774.1
```

```
sd(df$inventory)
```

```
#Returns: 1573358
```

```
plotrix::std.error(df$inventory)
#Returns: 15924.22
```

```
Hmisc::describe(df)
```

```
#VISUALIZING VARIABLE OF INTEREST#####
```

```
#Histogram of inventory, excluding values above 100K
ggplot(df,aes(x=inventory)) + geom_histogram(binwidth=1000,aes(fill=..count..)) +
  xlab('Inventory Value') + ylab('Count') + ggtitle('Counts of Inventory Value:
Excluding Values > 100K') +
  xlim(c(0, 100000)) + ylim(c(0,450))
```

```
#Histogram of inventory, excluding values above 1M
ggplot(df,aes(x=inventory)) + geom_histogram(binwidth=1000,fill='blue',alpha=0.6) +
  xlab('Inventory Value') + ylab('Count') + ggtitle('Counts of Inventory Value:
Excluding Values > 1M') +
  xlim(c(0, 1000000)) + ylim(c(0,450))
```

```
#####
#####
```

```
#Using highest valued Inventory records as the certainty stratum
# CF. https://enablement.acl.com/helpdocs/analytics/13/user-guide/en-us/Content/da\_sampling\_data/classical\_variables\_sampling.htm#navlink-5
```

```
#NB: The stratification::strata.LH() function below accepts an index vector as input
to the argument specifying the certainty stratum
cert_vec <- list(top_certainty_vec <- c(which(df$inventory==max(df$inventory))),
  top5_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:5],]$coID,
```

```

top10_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:10],]$coID,
top20_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:20],]$coID,
top50_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:50],]$coID,
top80_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:80],]$coID,
top90_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:90],]$coID,
top100_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:100],]$coID,
top110_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:110],]$coID,
top120_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:120],]$coID,
top150_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:150],]$coID,
top200_certainty_vec <- df[order(df$inventory,
decreasing=T)[1:200],]$coID)

```

#NB: Must define cert_vec in global environment before passing to minCV() function below. For common scoping issues:

```

# CF. https://stackoverflow.com/questions/28601959/error-object-not-found-in-nested-functions

```

```

#Generalized Lavallee-Hidiroglou Method -- applying Kozak's (2004)
Algorithm#####
# CF. https://www.rdocumentation.org/packages/stratification/versions/2.2-5/topics/strata.LH

```

```

#Function Optimizing the number of sampled strata and the best performing certainty
stratum from the vector list above#####

```

#NB: Passing the vector `c(0.5, 0, 0.5)` to the 'alloc' argument of the `stratification::strata.LH()` function applies Neyman allocation

Passing 'none' to the 'model' argument assumes that $Y=X$, i.e. that the variable being stratified is also the one being analyzed

```
minCV <- function(lo_Ls, hi_Ls, cert_vec){
  CV.list<-numeric()
  min.CV <- 999
  min.i <- 99
  min.j <- 99

  for (i in lo_Ls:hi_Ls){
    for (j in 1:length(cert_vec)){
      k <- tryCatch(stratification::strata.LH(x=df$`inventory`, n=500, Ls=i,
alloc=c(0.5,0,0.5),
                                certain=cert_vec[j], model="none", algo="Kozak")$RRMSE,
warning=function(w){
  message("Warning: LH Algo failed. Some sampled strata < minimum
Nh")
                                NA})
      if(is.na(k)){
        next
      }
      cat('Coefficient of Variation (RRMSE): ', k, '\n')
      cat('Selected Number of Strata: ', i, '\n')
      cat('Index of Certainty Stratum: ', j, '\n')

      CV.list <- append(CV.list, k)
      cat('All Coefficients of Variation: \n', CV.list, '\n')

      min_val <- min(CV.list, na.rm=TRUE)
      if (min.CV > min_val){
        min.CV <- min_val
        min.i <- i
        min.j <- j
      }
    }
  }
  result_list <- list('Minimum_CV'=min.CV, 'Best_LS'=min.i, 'Best_CertStratum'=min.j)
  return(result_list)
}
```

```

#Implement Stratification
Optimization#####
#####

best_results <- minCV(5, 10, cert_vec)
best_results$Minimum_CV
best_results$Best_LS
best_results$Best_CertStratum


#Best Stratification Configuration: Lavallee-Hidioglou (1988) Method and Kozak's
(2004) Algorithm#####

Optim_Config <- stratification::strata.LH(x=df$`inventory`, n=500, Ls=10,
alloc=c(0.5,0,0.5),
          certain=cert_vec[9], model='none', algo='Kozak')

print(Optim_Config)
plot(Optim_Config)


#Performance Comparison: Gunning and Horgan (2004) Geometric Stratification Method

#####
#####

Optim_Geo <- stratification::strata.geo(x=df$`inventory`, n=500, Ls = 8, alloc =
c(0.5, 0, 0.5),
          certain=cert_vec[9], model="none")

```



```
print(Optim_Geo)
plot(Optim_Geo)
```

```
#####
#####
```

```
#Apply Optimum Stratification to Sampling and
Estimation#####
```

```
#Preparing the data
frame#####
#####
```

```
#Identify boundary separating highest-valued certainty stratum from other strata
options(digits=15)
min(df[order(df$inventory, decreasing=T)[1:110],]$inventory)
#Returns 1372514.035
```

```
#Assign Strata
df$Stratum <- NA
df$Stratum[df$`inventory` < 36703.82] <- 1
df$Stratum[(df$`inventory` < 71655.62) & (df$`inventory` >= 36703.82)] <- 2
df$Stratum[(df$`inventory` < 106848.40) & (df$`inventory` >= 71655.62)] <- 3
df$Stratum[(df$`inventory` < 151104.80) & (df$`inventory` >= 106848.40)] <- 4
df$Stratum[(df$`inventory` < 209954.79) & (df$`inventory` >= 151104.80)] <- 5
df$Stratum[(df$`inventory` < 295186.09) & (df$`inventory` >= 209954.79)] <- 6
df$Stratum[(df$`inventory` < 413551.11) & (df$`inventory` >= 295186.09)] <- 7
df$Stratum[(df$`inventory` < 614205.40) & (df$`inventory` >= 413551.11)] <- 8
df$Stratum[(df$`inventory` < 974662.93) & (df$`inventory` >= 614205.40)] <- 9
df$Stratum[(df$`inventory` < 1372514.04) & (df$`inventory` >= 974662.93)] <- 10
df$Stratum[(df$`inventory` < 105379553.91) & (df$`inventory` >= 1372514.04)] <- 11
```

```
#Check if any NAs in Stratum column
any(is.na(df$Stratum))
```

```
#Assign column for Finite Population Correction
```

```
df$Num <- NA
```

```
df$Num[df$Stratum == 1] <- 4346
```

```
df$Num[df$Stratum == 2] <- 1047
```

```
df$Num[df$Stratum == 3] <- 1341
```

```
df$Num[df$Stratum == 4] <- 851
```

```
df$Num[df$Stratum == 5] <- 723
```

```
df$Num[df$Stratum == 6] <- 424
```

```
df$Num[df$Stratum == 7] <- 361
```

```
df$Num[df$Stratum == 8] <- 248
```

```
df$Num[df$Stratum == 9] <- 176
```

```
df$Num[df$Stratum == 10] <- 135
```

```
df$Num[df$Stratum == 11] <- 110
```

```
#Check if any NAs in df
```

```
any(is.na(df))
```

```
#Specify vector of population stratum sizes, appending certainty stratum
```

```
Nh <- append(Optim_Config$Nh, 110)
```

```
#Specify vector of population stratum standard deviations, appending 1 for certainty  
stratum
```

```
Sh <- append(sqrt(Optim_Config$varh), 1)
```

```
#NB: Appending 1 as the certainty stratum standard deviation may result in an  
underestimation of variance for this severely right-skewed population.
```

```
# See results of sampling procedures #1-10 below, which serve as a performance  
comparison, assigning only 1 or 2 units to the 'certainty' stratum.
```

```
# The results of procedures #12-16 accept the stratum sample sizes determined  
by LH-Optimization, assigning all 110 highest-valued records to the 'certainty'  
stratum.
```

```
#Neyman Allocation to determine strata sizes in sample of size n = 500
```

```
PracTools::strAlloc(n.tot=500, Nh=Nh, Sh=Sh, alloc='neyman')
```

```
#Returns sample stratum sizes of (100, 38, 43, 39, 41, 39, 43, 50, 60, 47, 0)
```

```
#Run the allocation again for a sample of size 499, reserving 1 unit for the
certainty stratum
```

```
Nh <- Optim_Config$Nh
```

```
Sh <- sqrt(Optim_Config$varh)
```

```
PracTools::strAlloc(n.tot=499, Nh=Nh, Sh=Sh, alloc='neyman')
```

```
#Returns sample stratum sizes of (101, 37, 43, 39, 41, 38, 43, 50, 60, 47)
```

```
#Adjusted Neyman Allocation, including certainty stratum
```

```
Ney_Alloc <- c(101, 37, 43, 39, 41, 38, 43, 50, 60, 47, 1)
```

```
#####
#####
```

```
#PASS 1 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- IMPLEMENTING LONELY
PSU ADJUSTMENT
```

```
#####
#####
```

```
#Set Random Sampling Seed
```

```
set.seed(222)
```

```
#Sampling using Neyman Allocation
```

```
Ney_Samp_1 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc,
method="srswor")
```

```
#Drop duplicated column
```

```
Ney_Samp_1 <- Ney_Samp_1[, !duplicated(colnames(Ney_Samp_1))]
```

```

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_1)[2] <- "coID"
Ney_Samp_1_merged <- merge(x = df, y = Ney_Samp_1)

#sort by Stratum
Ney_Samp_1_merged <- Ney_Samp_1_merged[order(Ney_Samp_1_merged$Stratum),]

#Neyman stratified design object
mydesign_1 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_1_merged,
fpc = ~Num)

#Adjust for single-PSU certainty stratum
options(survey.lonely.psu = "certainty")

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_1)
#Returns:              mean              SE
#      inventory 169756.4553944  3452.97697

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_1))
#Returns:              2.5 %              97.5 %
#      inventory 162988.744885283  176524.165903571

#DOES NOT CONTAIN TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_1)
#Returns:              total              SE
#      inventory  1657162517.56  33707961.22347

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_1))
#Returns:              2.5 %              97.5 %
#      inventory  1591096127.57013  1723228907.55066

```

```
#DOES NOT CONTAIN TRUE VALUE of 1754954822.94687#####
```

```
#PASS 2 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- IMPLEMENTING LONELY  
PSU ADJUSTMENT
```

```
#####  
#####
```

```
#Set Random Sampling Seed  
set.seed(223)
```

```
#Sampling using Neyman Allocation  
Ney_Samp_2 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc,  
method="srswor")
```

```
#Drop duplicated column  
Ney_Samp_2 <- Ney_Samp_2[, !duplicated(colnames(Ney_Samp_2))]
```

```
#Rename ID column to match original data frame, for merging  
colnames(Ney_Samp_2)[2] <- "coID"  
Ney_Samp_2_merged <- merge(x = df, y = Ney_Samp_2)
```

```
#sort by Stratum  
Ney_Samp_2_merged <- Ney_Samp_2_merged[order(Ney_Samp_2_merged$Stratum),]
```

```
#Neyman stratified design object  
mydesign_2 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_2_merged,  
fpc = ~Num)
```

```
#Adjust for single-PSU certainty stratum
options(survey.lonely.psu = "certainty")
```

```
#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_2)
#Returns:          mean          SE
#      inventory 171029.233252  3328.92537
```

```
#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_2))
#Returns:          2.5 %          97.5 %
#      inventory 164504.659423391 177553.807080663
```

```
#DOES NOT CONTAIN TRUE VALUE of 179774.106017914#####
```

```
#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_2)
#Returns:          total          SE
#      inventory 1669587375.006  32496969.44308
```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_2))
#Returns:          2.5 %          97.5 %
#      inventory 1605894485.29115 1733280264.72143
```

```
#DOES NOT CONTAIN TRUE VALUE of 1754954822.94687#####
```

```
#PASS 3 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- IMPLEMENTING LONELY
PSU ADJUSTMENT
```

```
#####  
#####
```

```
#Set Random Sampling Seed
```

```
set.seed(224)
```

```
#Sampling using Neyman Allocation
```

```
Ney_Samp_3 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc,  
method="srswor")
```

```
#Drop duplicated column
```

```
Ney_Samp_3 <- Ney_Samp_3[, !duplicated(colnames(Ney_Samp_3))]
```

```
#Rename ID column to match original data frame, for merging
```

```
colnames(Ney_Samp_3)[2] <- "coID"
```

```
Ney_Samp_3_merged <- merge(x = df, y = Ney_Samp_3)
```

```
#sort by Stratum
```

```
Ney_Samp_3_merged <- Ney_Samp_3_merged[order(Ney_Samp_3_merged$Stratum),]
```

```
#Neyman stratified design object
```

```
mydesign_3 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_3_merged,  
fpc = ~Num)
```

```
#Adjust for single-PSU certainty stratum
```

```
options(survey.lonely.psu = "certainty")
```

```
#Neyman stratified mean estimate and SE
```

```
survey::svymean(~`inventory`, design = mydesign_3)
```

```
#Returns:          mean          SE
```

```
#      inventory 170941.7207444 3432.95716
```

```
#Confidence interval of the mean estimate
```

```
confint(survey::svymean(~`inventory`, design = mydesign_3))
```

```
#Returns:          2.5 %          97.5 %
#      inventory 164213.248356339 177670.193132409
```

```
#DOES NOT CONTAIN TRUE VALUE of 179774.106017914#####
```

```
#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_3)
#Returns:          total          SE
#      inventory 1668733077.907 33512527.76587
```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_3))
#Returns:          2.5 %          97.5 %
#      inventory 1603049730.45458 1734416425.35857
```

```
#DOES NOT CONTAIN TRUE VALUE of 1754954822.94687#####
```

```
#PASS 4 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- IMPLEMENTING LONELY
PSU ADJUSTMENT
```

```
#####
#####
```

```
#Set Random Sampling Seed
set.seed(225)
```

```
#Sampling using Neyman Allocation
Ney_Samp_4 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc,
method="srswor")
```



```

#Drop duplicated column
Ney_Samp_4 <- Ney_Samp_4[, !duplicated(colnames(Ney_Samp_4))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_4)[2] <- "coID"
Ney_Samp_4_merged <- merge(x = df, y = Ney_Samp_4)

#sort by Stratum
Ney_Samp_4_merged <- Ney_Samp_4_merged[order(Ney_Samp_4_merged$Stratum),]

#Neyman stratified design object
mydesign_4 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_4_merged,
fpc = ~Num)

#Adjust for single-PSU certainty stratum
options(survey.lonely.psu = "certainty")

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_4)
#Returns:           mean           SE
#      inventory  180661.3660204  4638.56949

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_4))
#Returns:           2.5 %           97.5 %
#      inventory  171569.936881626  189752.795159095

#CONTAINS TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_4)
#Returns:           total           SE
#      inventory  1763616255.091  45281715.35415

```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_4))
#Returns:                2.5 %                97.5 %
#      inventory 1674865723.83843 1852366786.34309
```

```
#CONTAINS TRUE VALUE of 1754954822.94687#####
```

```
#PASS 5 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- IMPLEMENTING LONELY
PSU ADJUSTMENT
```

```
#####
#####
```

```
#Set Random Sampling Seed
set.seed(226)
```

```
#Sampling using Neyman Allocation
Ney_Samp_5 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc,
method="srswor")
```

```
#Drop duplicated column
Ney_Samp_5 <- Ney_Samp_5[, !duplicated(colnames(Ney_Samp_5))]
```

```
#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_5)[2] <- "coID"
Ney_Samp_5_merged <- merge(x = df, y = Ney_Samp_5)
```

```
#sort by Stratum
Ney_Samp_5_merged <- Ney_Samp_5_merged[order(Ney_Samp_5_merged$Stratum),]
```

```

#Neyman stratified design object
mydesign_5 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_5_merged,
fpc = ~Num)

#Adjust for single-PSU certainty stratum
options(survey.lonely.psu = "certainty")

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_5)
#Returns:          mean          SE
#      inventory 188972.1127207 4634.39056

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_5))
#Returns:          2.5 %          97.5 %
# q      inventory 179888.874126264 198055.351315211

#DOES NOT CONTAIN TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_5)
#Returns:          total          SE
#      inventory 1844745764.38 45240920.6795

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_5))
#Returns:          2.5 %          97.5 %
#      inventory 1756075189.22059 1933416339.53909

#DOES NOT CONTAIN TRUE VALUE of 1754954822.94687#####

```

```
#####  
#####
```

```
#PASS 6 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- NO LONELY PSU  
CERTAINTY STRATUM ADJUSTMENT  
#####  
#####
```

```
#Run the allocation again for a sample of size 498, reserving 2 units for the  
certainty stratum
```

```
Nh <- Optim_Config$Nh  
Sh <- sqrt(Optim_Config$varh)  
PracTools::strAlloc(n.tot=498, Nh=Nh, Sh=Sh, alloc='neyman')  
#returns sample stratum sizes of (101, 37, 43, 39, 41, 38, 42, 50, 60, 47)
```

```
#adjusted Neyman Allocation, including certainty stratum  
Ney_Alloc_2 <- c(101, 37, 43, 39, 41, 38, 43, 50, 60, 47, 2)
```

```
#Set Random Sampling Seed  
set.seed(227)
```

```
#Sampling using Neyman Allocation  
Ney_Samp_6 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_2,  
method="srswor")
```

```
#Drop duplicated column  
Ney_Samp_6 <- Ney_Samp_6[, !duplicated(colnames(Ney_Samp_6))]
```

```

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_6)[2] <- "coID"
Ney_Samp_6_merged <- merge(x = df, y = Ney_Samp_6)

#sort by Stratum
Ney_Samp_6_merged <- Ney_Samp_6_merged[order(Ney_Samp_6_merged$Stratum),]

#Neyman stratified design object
mydesign_6 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_6_merged,
fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_6)
#Returns:              mean              SE
#      inventory 185135.8284971 4675.96438

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_6))
#Returns:              2.5 %              97.5 %
#      inventory 175971.106721974 194300.550272168

#CONTAINS TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_6)
#Returns:              total              SE
#      inventory 1807295957.788 45646764.26414

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_6))
#Returns:              2.5 %              97.5 %
#      inventory 1717829943.81991 1896761971.7569

#CONTAINS TRUE VALUE of 1754954822.94687#####

```

```

#PASS 7 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- NO LONELY PSU
CERTAINTY STRATUM ADJUSTMENT
#####
#####

#Set Random Sampling Seed
set.seed(228)

#Sampling using Neyman Allocation
Ney_Samp_7 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_2,
method="srswor")

#Drop duplicated column
Ney_Samp_7 <- Ney_Samp_7[, !duplicated(colnames(Ney_Samp_7))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_7)[2] <- "coID"
Ney_Samp_7_merged <- merge(x = df, y = Ney_Samp_7)

#sort by Stratum
Ney_Samp_7_merged <- Ney_Samp_7_merged[order(Ney_Samp_7_merged$Stratum),]

#Neyman stratified design object
mydesign_7 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_7_merged,
fpc = ~Num)

#Neyman stratified mean estimate and SE

```

```

survey::svymean(~`inventory`, design = mydesign_7)
#Returns:          mean          SE
#      inventory  179848.508501    4643.21721

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_7))
#Returns:          2.5 %          97.5 %
#      inventory  170747.969987538  188949.047014456

```

#CONTAINS TRUE VALUE of 179774.106017914#####

```

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_7)
#Returns:          total          SE
#      inventory  1755681139.987  45327086.45115

```

```

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_7))
#Returns:          2.5 %          97.5 %
#      inventory  1666841683.01835  1844520596.95512

```

#CONTAINS TRUE VALUE of 1754954822.94687#####

#PASS 8 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- NO LONELY PSU
CERTAINTY STRATUM ADJUSTMENT

```

#####
#####

```

```

#Set Random Sampling Seed
set.seed(229)

#Sampling using Neyman Allocation
Ney_Samp_8 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_2,
method="srswor")

#Drop duplicated column
Ney_Samp_8 <- Ney_Samp_8[, !duplicated(colnames(Ney_Samp_8))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_8)[2] <- "coID"
Ney_Samp_8_merged <- merge(x = df, y = Ney_Samp_8)

#sort by Stratum
Ney_Samp_8_merged <- Ney_Samp_8_merged[order(Ney_Samp_8_merged$Stratum),]

#Neyman stratified design object
mydesign_8 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_8_merged,
fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_8)
#Returns:
              mean          SE
#      inventory 184075.5593905 4733.82151

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_8))
#Returns:
              2.5 %          97.5 %
#      inventory 174797.439716593 193353.679064313

#CONTAINS TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE

```



```

survey::svytotal(~`inventory`, design = mydesign_8)
#Returns:                total                SE
#      inventory    1796945610.77    46211565.60562

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_8))
#Returns:                2.5 %                97.5 %
#      inventory    1706372606.51338    1887518615.02582

#CONTAINS TRUE VALUE of 1754954822.94687#####

#PASS 9 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- NO LONELY PSU
CERTAINTY STRATUM ADJUSTMENT
#####
#####

#Set Random Sampling Seed
set.seed(230)

#Sampling using Neyman Allocation
Ney_Samp_9 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_2,
method="srswor")

#Drop duplicated column
Ney_Samp_9 <- Ney_Samp_9[, !duplicated(colnames(Ney_Samp_9))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_9)[2] <- "coID"

```

```

Ney_Samp_9_merged <- merge(x = df, y = Ney_Samp_9)

#sort by Stratum
Ney_Samp_9_merged <- Ney_Samp_9_merged[order(Ney_Samp_9_merged$Stratum),]

#Neyman stratified design object
mydesign_9 <- survey::svydesign(id = ~1, strata = ~Stratum, data = Ney_Samp_9_merged,
fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_9)
#Returns:              mean              SE
#      inventory  182583.5832636  5040.97741

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_9))
#Returns:              2.5 %              97.5 %
#      inventory  172703.449097735  192463.717429423

#CONTAINS TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_9)
#Returns:              total              SE
#      inventory  1782380939.819  49210021.45333

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_9))
#Returns:              2.5 %              97.5 %
#      inventory  1685931070.09209  1878830809.54603

#CONTAINS TRUE VALUE of 1754954822.94687#####

```

```

#PASS 10 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- NO LONELY PSU
CERTAINTY STRATUM ADJUSTMENT
#####
#####

#Set Random Sampling Seed
set.seed(231)

#Sampling using Neyman Allocation
Ney_Samp_10 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_2,
method="srswor")

#Drop duplicated column
Ney_Samp_10 <- Ney_Samp_10[, !duplicated(colnames(Ney_Samp_10))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_10)[2] <- "coID"
Ney_Samp_10_merged <- merge(x = df, y = Ney_Samp_10)

#sort by Stratum
Ney_Samp_10_merged <- Ney_Samp_10_merged[order(Ney_Samp_10_merged$Stratum),]

#Neyman stratified design object
mydesign_10 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
Ney_Samp_10_merged, fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_10)
#Returns:                mean                SE

```

```

#          inventory  170758.0515275  3455.03629

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_10))
#Returns:          2.5 %          97.5 %
#          inventory  163986.304831325  177529.798223672

#DOES NOT CONTAIN TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_10)
#Returns:          total          SE
#          inventory  1666940099.011  33728064.2754

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_10))
#Returns:          2.5 %          97.5 %
#          inventory  1600834307.7634  1733045890.25948

#DOES NOT CONTAIN TRUE VALUE of 1754954822.94687#####

```

```

#####
#####

```

```
#PASS 11: PERFORMANCE COMPARISON with UNSTRATIFIED SIMPLE RANDOM SAMPLING
#####
#####
```

```
#Set Random Sampling Seed
set.seed(1)
```

```
#SRS of Size 500
srs_df <- df[sample(nrow(df), 500), ]
```

```
#SRS design object
mydesign_srs <- survey::svydesign(id = ~1, data = srs_df)
```

```
#SRS mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_srs)
#Returns:
      mean      SE
#   inventory 314438.5186579 203292.38226
```

```
#confidence interval of SRS mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_srs))
#Returns:
      2.5 %      97.5 %
#   inventory -84007.2288961298 712884.266211934
```

```
#HUGE UNINFORMATIVE CONFIDENCE INTERVAL CONTAINS TRUE VALUE of 179774.106017914###
```

```
#SRS estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_srs)
#Returns:
      total      SE
#   inventory 157219259.329 101646191.1282
```

```
#Confidence interval of SRS estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_srs))
#Returns:
      2.5 %      97.5 %
#   inventory -42003614.4480649 356442133.105967
```

```
#HUGE UNINFORMATIVE CONFIDENCE INTERVAL STILL DOES NOT CONTAIN TRUE VALUE of  
1754954822.94687###
```

```
#####  
#####
```

```
#PASS 12 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- ASSUMING SAMPLE  
STRATUM SIZES ASSIGNED BY LH OPTIMIZATION  
#####  
#####
```

```
#adjusted Neyman Allocation, assuming values provided by the LH optimization  
Ney_Alloc_3 <- append(Optim_Config$nh, 110)
```

```
#Set Random Sampling Seed  
set.seed(232)
```

```
#Sampling using Neyman Allocation  
Ney_Samp_12 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,  
method="srswor")
```

```
#Drop duplicated column  
Ney_Samp_12 <- Ney_Samp_12[, !duplicated(colnames(Ney_Samp_12))]
```

```
#Rename ID column to match original data frame, for merging
```

```

colnames(Ney_Samp_12)[2] <- "coID"
Ney_Samp_12_merged <- merge(x = df, y = Ney_Samp_12)

#sort by Stratum
Ney_Samp_12_merged <- Ney_Samp_12_merged[order(Ney_Samp_12_merged$Stratum),]

#Neyman stratified design object
mydesign_12 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
Ney_Samp_12_merged, fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_12)
#Returns:              mean              SE
#      inventory 173299.0523965  7965.95442

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_12))
#Returns:              2.5 %              97.5 %
#      inventory 157686.068638111 188912.036154827

#CONTAINS TRUE VALUE of 179774.106017914#####

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_12)
#Returns:              total              SE
#      inventory 1691745349.494  77763647.01153

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_12))
#Returns:              2.5 %              97.5 %
#      inventory 1539331402.04524 1844159296.94342

#CONTAINS TRUE VALUE of 1754954822.94687#####

```

```

#PASS 13 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- ASSUMING SAMPLE
STRATUM SIZES ASSIGNED BY LH OPTIMIZATION
#####
#####

#adjusted Neyman Allocation, assuming values provided by the LH optimization
Ney_Alloc_3 <- append(Optim_Config$nh, 110)

#Set Random Sampling Seed
set.seed(233)

#Sampling using Neyman Allocation
Ney_Samp_13 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,
method="srswor")

#Drop duplicated column
Ney_Samp_13 <- Ney_Samp_13[, !duplicated(colnames(Ney_Samp_13))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_13)[2] <- "coID"
Ney_Samp_13_merged <- merge(x = df, y = Ney_Samp_13)

#sort by Stratum
Ney_Samp_13_merged <- Ney_Samp_13_merged[order(Ney_Samp_13_merged$Stratum),]

#Neyman stratified design object

```



```
mydesign_13 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
Ney_Samp_13_merged, fpc = ~Num)
```

```
#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_13)
#Returns:          mean          SE
#      inventory 170343.7363795   7660.27815
```

```
#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_13))
#Returns:          2.5 %          97.5 %
#      inventory 155329.867097901 185357.605661186
```

```
#CONTAINS TRUE VALUE of 179774.106017914#####
```

```
#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_13)
#Returns:          total          SE
#      inventory 1662895554.537   74779635.28079
```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_13))
#Returns:          2.5 %          97.5 %
#      inventory 1516330162.60971 1809460946.4645
```

```
#CONTAINS TRUE VALUE of 1754954822.94687#####
```

```

#PASS 14 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- ASSUMING SAMPLE
STRATUM SIZES ASSIGNED BY LH OPTIMIZATION
#####
#####

#adjusted Neyman Allocation, assuming values provided by the LH optimization
Ney_Alloc_3 <- append(Optim_Config$nh, 110)

#Set Random Sampling Seed
set.seed(234)

#Sampling using Neyman Allocation
Ney_Samp_14 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,
method="srswor")

#Drop duplicated column
Ney_Samp_14 <- Ney_Samp_14[, !duplicated(colnames(Ney_Samp_14))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_14)[2] <- "coID"
Ney_Samp_14_merged <- merge(x = df, y = Ney_Samp_14)

#sort by Stratum
Ney_Samp_14_merged <- Ney_Samp_14_merged[order(Ney_Samp_14_merged$Stratum),]

#Neyman stratified design object
mydesign_14 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
Ney_Samp_14_merged, fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_14)
#Returns:          mean          SE
#      inventory  175205.17016   7810.02253

```

```
#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_14))
#Returns:                2.5 %                97.5 %
#      inventory 159897.807284972 190512.533034961
```

```
#CONTAINS TRUE VALUE of 179774.106017914#####
```

```
#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_14)
#Returns:                total                SE
#      inventory 1710352871.102 76241439.92664
```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_14))
#Returns:                2.5 %                97.5 %
#      inventory 1560922394.7159 1859783347.48729
```

```
#CONTAINS TRUE VALUE of 1754954822.94687#####
```

```
#PASS 15 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- ASSUMING SAMPLE
STRATUM SIZES ASSIGNED BY LH OPTIMIZATION
#####
#####
```

```
#adjusted Neyman Allocation, assuming values provided by the LH optimization
```

```

Ney_Alloc_3 <- append(Optim_Config$nh, 110)

#Set Random Sampling Seed
set.seed(235)

#Sampling using Neyman Allocation
Ney_Samp_15 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,
method="srswor")

#Drop duplicated column
Ney_Samp_15 <- Ney_Samp_15[, !duplicated(colnames(Ney_Samp_15))]

#Rename ID column to match original data frame, for merging
colnames(Ney_Samp_15)[2] <- "coID"
Ney_Samp_15_merged <- merge(x = df, y = Ney_Samp_15)

#sort by Stratum
Ney_Samp_15_merged <- Ney_Samp_15_merged[order(Ney_Samp_15_merged$Stratum),]

#Neyman stratified design object
mydesign_15 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
Ney_Samp_15_merged, fpc = ~Num)

#Neyman stratified mean estimate and SE
survey::svymean(~`inventory`, design = mydesign_15)
#Returns:
              mean              SE
#      inventory 174342.7305529  8034.61217

#Confidence interval of the mean estimate
confint(survey::svymean(~`inventory`, design = mydesign_15))
#Returns:
              2.5 %              97.5 %
#      inventory 158595.180070016  190090.281035698

#CONTAINS TRUE VALUE of 179774.106017914#####

```

```

#Neyman stratified estimate of total and SE
survey::svytotal(~`inventory`, design = mydesign_15)
#Returns:                total                SE
#      inventory 1701933735.657    78433884.00301

#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_15))
#Returns:                2.5 %                97.5 %
#      inventory 1548206147.8435    1855661323.47048

```

```

#CONTAINS TRUE VALUE of 1754954822.94687#####

```

```

#PASS 16 OF STRATIFIED SAMPLING ESTIMATION of MEAN and TOTAL -- ASSUMING SAMPLE
STRATUM SIZES ASSIGNED BY LH OPTIMIZATION
#####
#####

```

```

#adjusted Neyman Allocation, assuming values provided by the LH optimization
Ney_Alloc_3 <- append(Optim_Config$nh, 110)

```

```

#Set Random Sampling Seed
set.seed(236)

```

```

#Sampling using Neyman Allocation

```

```
Ney_Samp_16 <- sampling::strata(df, stratanames='Stratum', size=Ney_Alloc_3,
method="srswor")
```

```
#Drop duplicated column
```

```
Ney_Samp_16 <- Ney_Samp_16[, !duplicated(colnames(Ney_Samp_16))]
```

```
#Rename ID column to match original data frame, for merging
```

```
colnames(Ney_Samp_16)[2] <- "coID"
```

```
Ney_Samp_16_merged <- merge(x = df, y = Ney_Samp_16)
```

```
#sort by Stratum
```

```
Ney_Samp_16_merged <- Ney_Samp_16_merged[order(Ney_Samp_16_merged$Stratum),]
```

```
#Neyman stratified design object
```

```
mydesign_16 <- survey::svydesign(id = ~1, strata = ~Stratum, data =
```

```
Ney_Samp_16_merged, fpc = ~Num)
```

```
#Neyman stratified mean estimate and SE
```

```
survey::svymean(~`inventory`, design = mydesign_16)
```

```
#Returns:                mean                SE
```

```
#      inventory  170959.9091492  7684.90266
```

```
#Confidence interval of the mean estimate
```

```
confint(survey::svymean(~`inventory`, design = mydesign_16))
```

```
#Returns:                2.5 %                97.5 %
```

```
#      inventory  155897.77670773  186022.04159064
```

```
#CONTAINS TRUE VALUE of 179774.106017914#####
```

```
#Neyman stratified estimate of total and SE
```

```
survey::svytotal(~`inventory`, design = mydesign_16)
```

```
#Returns:                total                SE
```

```
#      inventory  1668910633.114  75020019.78265
```

```
#Confidence interval of the estimate of total
confint(survey::svytotal(~`inventory`, design = mydesign_16))
#Returns:                2.5 %                97.5 %
#      inventory 1521874096.22086  1815947170.00782
```

```
#CONTAINS TRUE VALUE of 1754954822.94687#####
```

```
#####
#####
```

```
#SUMMARY OF
RESULTS#####
#####
```

```
#Mean of the Total and Mean Estimates and their Standard Errors (across Samples
assuming the Sizes allocated by LH Optimization)
```

```
m.est_of_mean <- mean(173299.05, 170343.74, 175205.17, 174342.73, 170959.91)
#Returns: 173299.05
```

```
m.est_of_mean_SE <- mean(7965.95, 7660.28, 7810.02, 8034.61, 7684.90)
#Returns: 7965.95
```

```
m.est_of_total <- mean(1691745349.49, 1662895554.54, 1710352871.10, 1701933735.66,
1668910633.11)
#Returns: 1691745349.49
```

```

m.est_of_total_SE <- mean(77763647.01, 74779635.28, 76241439.93, 78433884.00,
75020019.78)
#Returns: 77763647.01

```

```

#VISUALIZE TABLE OF
RESULTS#####
#####

```

```

val_df <- setNames(data.frame(matrix(ncol = 8, nrow = 16)),
c('Mean_Estimate', 'Mean_SE', 'Mean_95%_CI', 'Includes_True_Mean', 'Total_Estimate', 'Total_SE', 'Total_95%_CI', 'Includes_True_Total'))

```

```

rownames(val_df) <- c(paste('Stratified_Test', 1:5),
paste('Stratified_No_Lonely_PSU', 1:5), 'SRS', paste('Stratified_LH-Optimized', 1:5))

```

```

val_df[,1] <- c(169756.46, 171029.23, 170941.72, 180661.37, 188972.11,
185135.83, 179848.51, 184075.56, 182583.58, 170758.05,
314438.52,
173299.05, 170343.74, 175205.17, 174342.73, 170959.91)

```

```

val_df[,2] <- c(3452.98, 3328.93, 3432.96, 4638.57, 4634.39,
4675.96, 4643.22, 4733.82, 5040.98, 3455.04,
203292.38,
7965.95, 7660.28, 7810.02, 8034.61, 7684.90)

```

```

val_df[,3] <- c('[162988.75, 176524.17]', '[164504.66, 177553.81]', '[164213.25,
177670.19]', '[171569.94, 189752.80]', '[179888.87, 198055.35]',
'[175971.11, 194300.55]', '[170747.97, 188949.05]', '[174797.44,
193353.68]', '[172703.45, 192463.72]', '[163986.30, 177529.80]',
'[-84007.23, 712884.27]',

```



```
      '[157686.07, 188912.04]', '[155329.87, 185357.61]', '[159897.81, 190512.53]', '[158595.18, 190090.28]', '[155897.78, 186022.04]')
```

```
val_df[,4] <- c('NO', 'NO', 'NO', 'YES', 'NO',  
               'YES', 'YES', 'YES', 'YES', 'NO',  
               'YES*',  
               'YES', 'YES', 'YES', 'YES', 'YES')
```

```
val_df[,5] <- c(1657162517.56, 1669587375.01, 1668733077.91, 1763616255.09,  
1844745764.38,  
1807295957.79, 1755681139.99, 1796945610.77, 1782380939.82,  
1666940099.01,  
157219259.33,  
1691745349.49, 1662895554.54, 1710352871.10, 1701933735.66,  
1668910633.11)
```

```
val_df[,6] <- c(33707961.22, 32496969.44, 33512527.77, 45281715.35, 45240920.68,  
45646764.26, 45327086.45, 46211565.61, 49210021.45, 33728064.28,  
101646191.13,  
77763647.01, 74779635.28, 76241439.93, 78433884.00, 75020019.78)
```

```
val_df[,7] <- c('[1591096127.57, 1723228907.55]', '[1605894485.29, 1733280264.72]', '[1603049730.45, 1734416425.36]', '[1674865723.84, 1852366786.34]', '[1756075189.22, 1933416339.54]',  
               '[1717829943.82, 1896761971.76]', '[1666841683.02, 1844520596.96]', '[1706372606.51, 1887518615.03]', '[1685931070.09, 1878830809.55]', '[1600834307.76, 1733045890.26]',  
               '[-42003614.45, 356442133.11]',  
               '[1539331402.05, 1844159296.94]', '[1516330162.61, 1809460946.46]', '[1560922394.72, 1859783347.49]', '[1548206147.84, 1855661323.47]', '[1521874096.22, 1815947170.01]')
```

```
val_df[,8] <- c('NO', 'NO', 'NO', 'YES', 'NO',  
               'YES', 'YES', 'YES', 'YES', 'NO',  
               'NO',  
               'YES', 'YES', 'YES', 'YES', 'YES')
```

```
val_tbl <- tableGrob(val_df)
```

```
grid.arrange(val_tbl)
```